

Vision: Applying for a PhD program in the field of MLsys, especially about LLM Serving and Compilers.

EDUCATION

Bachelor of Computer Science Huazhong University of Science and Technology, GPA: **3.95/4.00** **Sept. 2020 — June 2024**

ACADEMIC EXPERIENCE

LLM Serving, Inferencing and Profiling University of California San Diego **Aug. 2023 — Present**

- **Advisor: Prof. Hao Zhang**
- Role: Research Intern
- Profiled the bottleneck of current SOTA LLM Serving framework(e.g., vllm, ppl.llm).
- Improve the GPU SM utilization to accelerate the serving throughput of LLMs.

WhiteFox: White-box Compiler Fuzzing via LLMs University of Illinois Urbana-Champaign **June. 2023 — Sept. 2023**

- **Advisor: Prof. Lingming Zhang**
- Role: Research Intern, **Third author**, paper already submitted to **FSE'24**
- Use LLMs to generate examples to trigger the optimization in the compiler
- Detect the flaws of Pytorch, XLA, LLVM based on the pattern written in the source code

Efficient Paged Dynamic Graph Serving HUST **Oct. 2022 — June 2023**

- **Advisor: Prof. Hai Jin, Prof. Zhiyuan Shao**
- Role: Research Intern, **Co-first author**, paper will be submitted to **ICDE'24**
- Remap the PageTable of OS Kernel and propose a new B+ Tree-based memory algorithm to accelerate the dynamic graph processing system.
- Speed up existing SOTA algorithms by more than 10x times.

INDUSTRIAL EXPERIENCE

Optimize the LLVM Backend of SenseTime TPU, GPU Compiler, Sensetime Company **April 2023 — Aug.2023**

- Role: LLVM Backend Developer *Shanghai. China*
- Mentor: Wenqiang Yin
- GPU Compiler Optimization, Instruction Selection, Instruction Pattern Match, CodeGen Emitter

Develop High Performance Neural Network Inference Engine, Tencent Company **July 2022 — Nov. 2022**

- Role: **Top 15** committer of 263(util Nov.2022) *Shenzhen. China*
- Mentor: nihui, with **6k+** followers in Github
- Optimize high performance operators and math library for ncnn, **18k+** stars in Github, a neural network library handcraftly optimized for X86/ARM/RISCV/GPU platforms.

Deploy High-FPS AI Models on Arm Chips, FiberHome Telecommunication Company **Dec. 2021 — June 2022**

- Role: **Leader** *Wuhan. China*
- Mentor: Yayu Gao, Xinggang Wang
- Deploy YOLOX/LiteHRNet on Snapdragon 870(Arm CPU), Achieve 20 FPS.

SKILLS

AI	LLM/CV Model Deployment
HPC	CUDA, Intel SSE, Arm NEON, Assembly
Compiler	Compiler Infra like LLVM, MLIR, Triton

MORE INFO

For better reading experience and more detailed information, please feel free to visit my  website :)